

Mohammed Siddiq

AI ML Engineer • MLOps & NLP • Enterprise Systems

MohammedSiddiq1235@gmail.com | 7981901203 | Hyderabad, India | [GitHub](#) | [LinkedIn](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

AI/ML Engineer with a B.Tech in AI & Machine Learning and AWS Cloud Practitioner certification, driving enterprise system reliability at Tech Mahindra for a US-based client. Skilled across the full ML lifecycle - NLP, transformer fine-tuning, RAG systems, and MLOps - backed by practical AWS knowledge in compute, storage, security, and well-architected cloud design. Independently reduced LMS engagement drops by 30% through content personalization and grew client acquisition 40% via a solo-built full-stack product. Seeking impactful roles at the intersection of AI research and production systems.

EXPERIENCE

Technical Operations Engineer | Tech Mahindra

Aug 2025 – Present • Hyderabad, India

US-Based Enterprise Client Engagement

- Triaged ~10–12 daily incidents across 2,000+ retail sites by querying structured databases with multi-table SQL joins and aggregate functions, analysing data from multiple sources to pinpoint root causes and dispatch repair services, directly reducing Time to Resolve (TTR) and improving SLA compliance.
- Coordinated with field technicians and lead specialists on escalated, complex issues to drive timely resolution and minimize site downtime across a 30+ member operations team.
- Analysed structured operational data across disparate sources to identify recurring failure patterns and systemic issues, enabling proactive maintenance interventions and reducing repeat incident frequency.
- Maintained accurate technical documentation, incident records, configuration changes, and resolution details to ensure traceability and consistent operational support.

Machine Learning Intern Ignitus

Feb 2024 – Jul 2024 • Remote

- Developed predictive modeling and NLP solutions using Python, Pandas, and Scikit-learn, applying feature engineering to improve model performance.
- Engineered end-to-end data pipelines and ETL workflows to clean, transform, and preprocess structured and unstructured data for machine learning applications.
- Optimized ML inference and data processing pipelines through code refactoring, improving execution efficiency and application performance.
- Collaborated in an Agile development environment using Git, feature branches, pull requests, and peer reviews to build and integrate 3+ reusable ML components for testing and deployment.

Data Analyst & Web Developer (Contract) | Excel Placement Services **India**

Oct 2024 – Mar 2025 • Hyderabad,

- Analyzed recruitment data using SQL, Excel, and Power BI, building KPI dashboards for Time-to-Hire, Submission-to-Interview, and recruitment pipeline performance.
- Automated reporting workflows using Excel, streamlining recruiter reporting processes and improving operational productivity.
- Designed, developed, and deployed a responsive full-stack recruitment website, managing the complete project lifecycle from requirements gathering and UI development to production deployment, contributing to a 40% increase in client acquisition inquiries.
- Led the end-to-end website migration by configuring GoDaddy hosting, DNS records, domain setup, backend components, and production deployment, while optimizing application latency from 1–1.5 seconds to under 500 ms with minimal downtime.

PROJECTS

Natural Language Question Answering System

- Built a Flask-based REST API enabling users to query unstructured PDF documents via transformer-powered NLP. Handled full ownership from model selection and API design to backend integration and deployment documentation.

Tech Stack: *Python • Flask • BERT • Hugging Face Transformers • HTML • CSS • JavaScript • REST API*

NASA Astronomy Picture of the Day (APOD)

- Developed a cloud-enabled web application that integrates with NASA's Astronomy Picture of the Day (APOD) API, allowing users to view daily images, search by date, and explore random astronomy content through a responsive interface.
- Implemented a serverless API proxy using AWS Lambda, Amazon API Gateway, CloudFront, and IAM to gain hands-on experience with cloud-native application architecture, API integration, and secure request routing while deploying the frontend on GitHub Pages.

Tech Stack: *HTML5, CSS3, JavaScript, REST API, AWS Lambda, Amazon API Gateway, Amazon CloudFront, AWS IAM, GitHub Pages*

RAG-Powered Multi-Document QA System

- Engineered an end-to-end Retrieval-Augmented Generation pipeline with multi-format document ingestion (PDF, DOCX, TXT), semantic vector storage via ChromaDB, and advanced retrieval strategies (MMR + MultiQueryRetriever), delivering cited, cross-document synthesised answers at scale using LLaMA 3.3 70B via Groq API with full source attribution and strategy-switching CLI.

Tech Stack: *Python • LangChain • ChromaDB • HuggingFace Sentence Transformers • Groq API • LLaMA 3.3 70B*

TECHNICAL SKILLS

Languages: Python, SQL

ML / AI: NLP, Deep Learning, Generative AI, LLMs, RAG Pipelines, Vector Embeddings, Prompt Engineering, Transfer Learning

Frameworks & Tools: Computer Vision, PyTorch, Keras, Hugging Face, LangChain, Pandas, NumPy, Flask

MLOps & Cloud: MLflow, DVC, Docker, GitHub Actions, AWS (EC2, S3, Lambda)

EDUCATION

B.Tech – Artificial Intelligence & Machine Learning

2020 – 2024

Siddhartha Institute of Technology and Science, JNTUH

- Relevant Coursework: Machine Learning, Deep Learning, Natural Language Processing, Data Structures & Algorithms, Database Management Systems, Computer Networks

CERTIFICATIONS

- **AWS** Cloud Practitioner Essentials – AWS Training & Certification (Jul 2026)
- Machine Learning Training – Internshala
- Deep Learning with TensorFlow – IBM
- Machine Learning with Python – freeCodeCamp
- Responsive Web Design – freeCodeCamp